

Niyaz

Software Engineer

Chicago | Open to Relocation | (940) 300 9580 | nniyaz0723@gmail.com

SUMMARY

Software Engineer in building **full-stack applications**, **AI-powered products**, and **distributed systems** for enterprise-scale platforms. Hands-on expertise across RAG pipelines, LLM integration, and event-driven microservices, working with **Apache Kafka**, **React**, **TypeScript**, **PostgreSQL**, and **AWS** to **design** and **ship cloud-native systems**. Improved AI answer accuracy from **71% to 92% (+21 points)**, reduced query p95 latency by 72%, and cut data freshness from **40 minutes to 6 minutes** (85% faster) across **high-scale production** applications.

TECHNICAL SKILLS

AI & ML Infra: ML Feature Stores, MLOps, LLM Integration Pipelines, Real-Time AI Data Pipelines, Recommendation Systems, Data Freshness Optimization, Apache Iceberg, Model Serving Infrastructure

Languages: Java, Go, Python, TypeScript, SQL, Node.js, Bash, C++, HTML/CSS

System Design: Distributed Systems, Microservices Architecture, Event-Driven Architecture, CQRS, Event Sourcing, Saga Pattern, Outbox Pattern, High Availability, Scalability, Software Design Architecture

Backend & APIs: Spring Boot, Gin (Go), gRPC + Protobuf, Federated GraphQL (DGS), REST APIs, WebSockets, OAuth 2.0, JWT, Spring Cloud Gateway, Swagger/OpenAPI

Streaming & Data: Apache Kafka, Apache Flink, Kafka Streams, Avro + Schema Registry, CDC (Debezium), Kafka Connect, Apache Iceberg, Cassandra, PostgreSQL, Redis, Elasticsearch, Snowflake, DynamoDB

Cloud & Platform: AWS (EKS, Kinesis, Lambda, RDS, SQS, S3, Secrets Manager), Terraform (IaC), Helm, Argo CD, GitOps, FinOps, Cost Optimization, Git, GitHub

Containers & SRE: Kubernetes, Docker, Istio Service Mesh, mTLS, HPA, Canary Deployments, Chaos Engineering, Resilience4j, OpenTelemetry, Prometheus, Grafana, Jaeger, ELK Stack, PagerDuty, SLO/SLI, MTTR/MTTD, Incident Response

Security: HashiCorp Vault, AWS Secrets Manager, Zero-Trust Architecture, mTLS, OAuth 2.0, JWT, Secrets Rotation Automation

Practices: Platform Engineering, A/B Testing, Agile/Scrum, Code Review, DevSecOps, Capacity Planning, On-Call Incident Response, SDLC, TDD

PROFESSIONAL EXPERIENCE

ServiceNow — Software Engineer | San Francisco

Dec 2025 - Present

- Reduced personalized **content retrieval latency** from 4.8s → 1.2s by building a real-time data platform using **Kafka**, **FastAPI**, and **React + TypeScript** in **Agile** sprints tracked via **Jira**, shipping through **Git-based CI/CD pipelines** and debugging data-processing failures across recommendation services.
- Boosted AI content-assistant accuracy from 72% → 93% (**RAGAS**) by engineering hybrid retrieval with **dense embeddings** and **cross-encoder reranking** on pgvector, validated through **pytest** and Jest with peer **code review** for the backend and frontend serving 50K+ MAUs.
- Optimized AI query** p95 latency from **3.6s → 850ms** by implementing **Redis distributed caching** and **SSE** streaming behind a **Zuul 2 gateway**, sustaining **600+ QPS** at peak while communicating with product and QA teams on **cross-functional** releases.
- Drove down failed inter-service requests from **2.4% → 0.6%** across 8 microservices by implementing circuit breakers and exponential-backoff retries with **mTLS-enforced** service communication, owning the observability rollout using **OpenTelemetry** tracing with **PagerDuty** on-call routing and **troubleshooting** dependency failures.

Cisco — Software Engineer | India

Jun 2022 – Aug 2024

- Scaled reconciliation throughput from **45K → 110K** accounts/hour across 1.2M active accounts by decomposing a legacy monolith into **9 microservices & Software Architecture Design** using **Kafka-based CQRS** and **AWS SQS** for async **event-driven processing** on **Linux-based AWS EKS clusters**.
- Trimmed** compliance document analysis from 14 min → 5 min per case and raised accuracy from **78% → 95%** by delivering an **LLM orchestration** and **pgvector** RAG workflow with citation grounding through a React annotation UI with **pytest** test coverage.
- Accelerated partner API configuration from **12 days → 4 days** by engineering a unified integration layer with **OAuth 2.0** and **Spring Cloud Gateway**, **delivering a React + TypeScript portal** with **GraphQL** adopted by 3 external banking partners.
- Lowered inter-service latency from **340ms → 130ms** by migrating a high-throughput settlement service from REST polling to **gRPC + Protocol Buffers streaming RPCs**, applying **OOP design** principles and efficient data structures to solve throughput **bottlenecks**.
- Compressed deployment recovery from 42 min → 9 min and raised successful daily releases from **92% → 99.5%** by deploying containerized services on **AWS EKS via Terraform IaC** and Argo CD with **Spinnaker** blue-green pipelines, mentoring 1 junior engineer on **GitOps** workflows across the **SDLC**.

EDUCATION

Master's in advanced data Analytics | University of North Texas. Texas

Bachelors in computer science | Gitam University, Visakhapatnam, India.